

Filtering before testing

Pre-selecting LD clusters as test units, and what it does and does not buy

module_sim_LDscnR

01 September 2026

Contents

1	Scope, and what changed after review	1
1.1	Terminology	2
2	Why the ranking approach was abandoned	2
3	Design	2
3.1	What is compared	2
3.2	Composition of the panels	2
3.3	The multiplicity mechanism, and its cost	3
4	Results, pooled — with ties and magnitudes	4
5	Stratified results	5
6	Null behaviour: two nulls, and they disagree	7
6.1	Which null is the right one	8
6.2	What this does and does not establish	9
7	Status of each claim	9
8	Relationship to the empirical panel analysis	9
8.1	What EMMAX is actually run on	9
8.2	What the representative costs, measured	10
8.3	Relationship to the empirical panel analysis	10
9	Corrections applied after review	11
9.1	Open, and in progress	11
10	Provenance	11

1 Scope, and what changed after review

This describes a **selection** analysis: choosing which units to test before the phenotype is seen. It shares a premise with the abandoned C-score work — ld_w is phenotype-blind and may carry information about where causal variation sits — but differs in unit, truth, comparison and conclusion.

An earlier draft was reviewed externally. That review found two factual errors, an ambiguous denominator, and — most consequentially — that pooling 80 panels across four parameter cells concealed the main structure in the data. **This version is stratified, and the conclusion changed as a result.** All corrections are recorded in Section~9 rather than silently applied.

The finding, as the evidence now supports it. Pre-filtering LD clusters increases QTN-proximal discoveries at moderate-to-large filter sizes — by a median of about **one** such discovery per panel, against four to five extra total calls. Which filter is better **depends on the regime**: local-LD weight outperforms cluster size where background LD is low, and fails where it is high. Neither filter is established as generally preferable.

1.1 Terminology

The endpoint is **not** a true discovery. Truth here is bp proximity to a driving QTN with the causal variants themselves excluded, so a cluster at 49 kb counts and one at 51 kb does not, regardless of actual LD. Throughout this document such discoveries are called **QTN-proximal**. “True discovery” is reserved for a rule tied to the simulated causal architecture, which this is not.

2 Why the ranking approach was abandoned

Recorded for self-containment. All paired within run with a sign test.

Test	Result
$\tau_C = 0.05$ vs $\alpha = 0.05$	C better in 79/151 non-tied (52%), $p = 0.63$
PR-AUC at the operative distance cap	C better in 120/313 (38%), $p = 4.4 \times 10^{-5}$
Stage-2 clusters as fixed units	α ahead by 0.07 PR, flat at every size floor

Two external results agree: the 3sp panel session finds C *ties* an oracle best-single-cell, and the simulation-parsing session measured the mechanism — reweighting a p-value by ld_w nets about **one marker in 1,400** at 50–100 kb and nothing at 10 kb.

3 Design

3.1 What is compared

Component	Choice	Why
Unit	Stage-2 LD cluster, tested through its LD-central pruned representative	Representative = highest median r^2 to the other members. Membership recomputed per bundle and required to reproduce the stored GRM marker set exactly
Filter	Top k units by median ld_w ; computed before any phenotype	Compared against MAF, cluster size and random at identical k
Test	BH at 0.05 within the selected set only	Identical to the conventional analysis, applied to fewer tests
Endpoint	QTN-proximal: within 10 / 50 / 100 kb of a driving QTN, causal variants excluded	Shares no machinery with ld_w ; an r^2 -tagging truth would flatter it by construction
Baseline	BH at 0.05 over all units	This IS the conventional α analysis, so the comparison is built in

3.2 Composition of the panels

The 80 panels are 4 parameter cells $\times$ 2 background-selection arms $\times$ 10 environments, each pooling 10 map sets into a genome-wide panel of about 86,000 stage-2 clusters. Cells and arms are **not** interchangeable and are stratified throughout Section~5.

3.3 The multiplicity mechanism, and its cost

Filtering changes no p-value. It changes the BH threshold, because BH over 99,786 units is far more severe than BH over 5,000. On the illustrative panel below the accounting is **27 discoveries** – **1 lost** + **28 gained** = **54**: filtering can and does *remove* an already-significant result when it falls outside the selected set. Of the 28 gained, **6 (21%) are QTN-proximal**, so precision on gained units is low by construction.




Figure 1: The mechanism. Navy: cleared the genome-wide cutoff. Orange: lie between the two cutoffs, gained by filtering alone. Grey dashed verticals mark driving QTN. One navy unit outside the selected set is lost and is not shown separately.




Figure 2: Units coloured by r^2 with the nearest driving QTN; triangles are the 25 clusters carrying a causal variant, which are excluded from testing. Only 459 of 99,786 clusters reach $r^2 \geq 0.2$ with any QTN, and several QTN-carrying clusters sit at $p \approx 1$.

4 Results, pooled — with ties and magnitudes

Filter	k	wins	losses	ties	% non-tied	median	IQR	p
ld_w	500	33	24	23	58	+0.0	[-2, 2]	0.289
ld_w	1,000	41	24	15	63	+1.0	[-1, 2]	0.046
ld_w	2,000	46	17	17	73	+1.0	[0, 3]	3e-04
ld_w	5,000	52	14	14	79	+1.0	[0, 3]	3e-06
ld_w	10,000	50	12	18	81	+1.0	[0, 2]	1e-06
ld_w	20,000	47	0	33	100	+1.0	[0, 2]	1e-14
size	500	50	17	13	75	+1.0	[0, 3]	7e-05
size	1,000	51	14	15	78	+1.0	[0, 3]	4e-06
size	2,000	49	11	20	82	+1.5	[0, 3]	8e-07
size	5,000	50	11	19	82	+1.0	[0, 3]	5e-07
size	10,000	46	10	24	82	+1.0	[0, 2]	1e-06
size	20,000	38	9	33	81	+0.0	[0, 1]	2e-05
MAF	500	4	49	27	8	-2.0	[-5, 0]	7e-11
MAF	1,000	7	46	27	13	-1.0	[-5, 0]	4e-08
MAF	2,000	10	46	24	18	-1.0	[-4, 0]	1e-06
MAF	5,000	15	43	22	26	-1.0	[-3, 0]	3e-04
MAF	10,000	17	41	22	29	-1.0	[-2, 0]	0.002
MAF	20,000	28	15	37	65	+0.0	[0, 1]	0.066
random	500	1	54	25	2	-2.0	[-5, 0]	3e-15
random	1,000	1	53	26	2	-2.0	[-5, 0]	6e-15
random	2,000	1	54	25	2	-2.0	[-5, 0]	3e-15
random	5,000	1	53	26	2	-2.0	[-4, 0]	6e-15
random	10,000	0	52	28	0	-2.0	[-4, 0]	4e-16
random	20,000	1	53	26	2	-1.0	[-3, 0]	6e-15

Read the denominator carefully. At $k = 20,000$, ld_w has 47 wins and 0 losses — 100% of non-tied panels — but 33 of 80 panels are **ties**, so the win is on 47 panels, not 80. The median gain is +1 QTN-proximal discovery.

Table 1: Median paired change against the alpha analysis: QTN-proximal discoveries, and total discoveries. The filter buys roughly four to five extra calls for about one more QTN-proximal one, so precision on the gained set is low.

k	QTN-proximal_MAF	QTN-proximal_LD	QTN-proximal_random	QTN-proximal_size	all discoveries_MAF	all discoveries_LD	all discoveries_random	all discoveries_size
500	-2	0	-2	1.0	-5	0	-5	0.0
1000	-1	1	-2	1.0	-5	1	-5	1.0
2000	-1	1	-2	1.5	-4	1	-5	1.0
5000	-1	1	-2	1.0	-3	4	-5	2.0
10000	-1	1	-2	1.0	-2	4	-4	1.5
20000	0	1	-1	0.0	0	5	-4	1.0

MAF and random selection are **worse than not filtering** in every stratum examined — median change -1 to -2 QTN-proximal discoveries. Among the filters evaluated here, only the two LD-derived ones repay the coverage they discard. Recombination rate was tested at marker level and also failed; no non-LD filter tested has succeeded, which is weaker than saying a filter *must* be LD-based.

5 Stratified results

Pooling concealed the main structure. Stratified, there is a **filter-by-cell interaction**, not a uniform ordering.

Cell	wins	losses	ties	% non-tied	median	p	Reading
V0.5_c1	18	0	2	100	+3.0	8e-06	ld_w much better
V2_c1	14	1	5	93	+1.0	1e-03	ld_w better
V1_c1.5	6	4	10	60	+0.0	0.754	no difference
V0.5_c2	1	18	1	5	-8.5	8e-05	cluster size much better

ld_w beats cluster size decisively in two cells, ties in one, and loses decisively in one. Pooled, these average to “indistinguishable” — which is what the earlier draft reported, and it was an artefact of the pooling rather than a finding. V0.5_c2 is the high-background-LD cell ($b = 0.42$); a plausible but untested reading is that ld_w is a *local*-LD statistic and loses discrimination when everything is correlated, whereas large clusters still mark real linkage blocks.




Figure 3: Filter-then-test against the conventional alpha analysis. Points are medians over panels; ribbons are the interquartile range across panels. The dashed line in panel A is the median genome-wide result; the dashed line in panel B is 50 percent of non-tied panels.

Against the alpha analysis, $k = 5,000$, 50 kb , by cell and BGS arm						
Filter	Cell / arm	wins	losses	ties	median	p
ld_w	V0.5_c1 bgs	10	0	0	+3.5	0.002
	V0.5_c1 nobgs	10	0	0	+4.0	0.002
	V0.5_c2 bgs	2	6	2	-3.5	0.289
	V0.5_c2 nobgs	3	7	0	-3.0	0.344
	V1_c1.5 bgs	3	1	6	+0.0	0.625
	V1_c1.5 nobgs	8	0	2	+1.0	0.008
	V2_c1 bgs	8	0	2	+1.0	0.008
	V2_c1 nobgs	8	0	2	+2.0	0.008
size	V0.5_c1 bgs	6	2	2	+1.0	0.289
	V0.5_c1 nobgs	7	3	0	+1.0	0.344
	V0.5_c2 bgs	9	0	1	+6.0	0.004
	V0.5_c2 nobgs	10	0	0	+3.5	0.002
	V1_c1.5 bgs	3	0	7	+0.0	0.250
	V1_c1.5 nobgs	6	2	2	+1.0	0.289
	V2_c1 bgs	3	2	5	+0.0	1.000
	V2_c1 nobgs	6	2	2	+1.0	0.289

Table 2: Background selection reduces the gain: the nobgs arm shows larger and more consistent improvement than the bgs arm at every k . Consistent with BGS costing recall rather than enrichment.

arm	k	wins	losses	ties	% non-tied	median	p
bgs	500	15	13	12	54	0.0	0.851
bgs	1000	17	12	11	59	0.0	0.458
bgs	2000	21	9	10	70	1.0	0.043
bgs	5000	23	7	10	77	1.0	0.005
bgs	10000	20	7	13	74	0.5	0.019
bgs	20000	20	0	20	100	0.5	2e-06
nobgs	500	18	11	11	62	0.0	0.265
nobgs	1000	24	12	4	67	1.0	0.065
nobgs	2000	25	8	7	76	1.5	0.005
nobgs	5000	29	7	4	81	2.0	3e-04
nobgs	10000	30	5	5	86	1.5	2e-05
nobgs	20000	27	0	13	100	1.0	1e-08

6 Null behaviour: two nulls, and they disagree

An earlier draft reported one null and dismissed the other. That was wrong, and the dismissed one carries the more important result.

Null	Selection	k	rejection rate	95% CI
genetic	unfiltered scan	all	0.042	[0.030, 0.056]
	ld_w	500	0.039	[0.028, 0.053]
	ld_w	1,000	0.033	[0.023, 0.046]
	ld_w	5,000	0.027	[0.018, 0.039]
	ld_w	20,000	0.031	[0.021, 0.044]
	ld_w	50,000	0.029	[0.020, 0.041]
	MAF	500	0.036	[0.025, 0.049]
	MAF	1,000	0.041	[0.030, 0.055]
	MAF	5,000	0.028	[0.019, 0.040]
	MAF	20,000	0.033	[0.023, 0.046]
	MAF	50,000	0.036	[0.025, 0.049]
	random	500	0.048	[0.036, 0.063]
	random	1,000	0.053	[0.040, 0.069]
	random	5,000	0.043	[0.031, 0.057]
	random	20,000	0.049	[0.036, 0.064]
	random	50,000	0.046	[0.034, 0.061]
env_orth	unfiltered scan	all	0.338	[0.309, 0.368]
	ld_w	500	0.073	[0.058, 0.091]
	ld_w	1,000	0.077	[0.061, 0.095]
	ld_w	5,000	0.067	[0.052, 0.084]
	ld_w	20,000	0.057	[0.043, 0.073]
	ld_w	50,000	0.147	[0.126, 0.170]
	MAF	500	0.047	[0.035, 0.062]
	MAF	1,000	0.059	[0.045, 0.075]
	MAF	5,000	0.075	[0.059, 0.093]
	MAF	20,000	0.078	[0.062, 0.096]
	MAF	50,000	0.080	[0.064, 0.099]
	random	500	0.094	[0.077, 0.114]
	random	1,000	0.104	[0.086, 0.125]
	random	5,000	0.124	[0.104, 0.146]
	random	20,000	0.180	[0.157, 0.205]
	random	50,000	0.259	[0.232, 0.287]

Both nulls are surrogate **environments** over 1,000 draws (10 environments $\times$ 100 surrogates), and under both $\lambda_{\text{sur}} \approx 0.99$, so neither result is genomic inflation.

The conventional scan fails the structure-preserving null. Under **env_orth** the unfiltered genome-wide analysis — the α baseline that every power comparison in this document is measured against — rejects in 34% of surrogates at $\alpha = 0.05$. Filtering reduces that to 0.057, an 83% reduction, **but does not repair it**: the residual still exceeds α .

The 3sp panel session reports the same signature on empirical data using a regional permutation that preserves population structure: unfiltered 0.37, filtered 0.21–0.34. Two datasets, two null constructions, one conclusion — *filtering improves calibration under structure and cannot rescue a scan that is already anti-conservative*.

6.1 Which null is the right one

genetic surrogates leave the scan calibrated at 0.042; **env_orth** surrogates do not. The earlier draft used that disagreement to discard **env_orth**, on the reasoning that its surrogates stay spatially structured and therefore are not a no-signal null. **That reasoning was convenient rather than correct.** The surrogates are orthogonal to the causal environment by construction, so associations recovered under them arise from shared spatial structure rather than from causation — which makes them false positives, and makes **env_orth**

a valid and arguably more realistic null. Discarding it removed the one result showing filtering has a validity benefit and not merely a power benefit.

Both are reported here. Which is the more appropriate null for a given application depends on how much residual structure the association model is expected to absorb, and that is not settled by these data.

6.2 What this does and does not establish

Every hypothesis is null in both experiments, so FDR equals the probability of at least one rejection and these are valid **global-null** checks. They do **not** establish FDR control under mixtures of null and alternative hypotheses, nor the Bourgon et al. independence condition in general.

They also do **not license choosing** k . They license evaluating a fixed, prespecified k under the tested null. Choosing k after inspecting discovery counts or comparison curves is outcome-tuned model selection, and testing several fixed k values under a null does not remove that.

Coverage. Both nulls use V2_c1, nobgs only, while the power experiment spans four cells and both arms. Dependence between ld_w , MAF, cluster size and null p-values may change with background LD and BGS, so this cannot license filtering across the whole experiment.

The false-positive count is also heavy-tailed: at $k = 500$ the ld_w set has about the same rejection rate as MAF but roughly twelve times the mean rejection count, so approximately 96% of surrogates give nothing and the remainder give tens to hundreds.

7 Status of each claim

Claim	Status	Basis
Filtering LD clusters increases QTN-proximal discoveries at moderate-to-large k	Supported	Pooled and in 3 of 4 cells; median gain +1, IQR [0, 3]
Among filters evaluated, only LD-derived ones repay the lost coverage	Supported	MAF, random and recombination all worse than no filtering
Which LD filter is better depends on regime	Established	Filter-by-cell interaction; ld_w wins 2 cells, loses 1 decisively
ld_w is generally preferable to cluster size	Not supported	Fails in V0.5_c2 by a median of -8.5
Global-null rejection rate is controlled under a genetic-basis null	Supported	One cell, one arm; CIs overlap α
The unfiltered scan is anti-conservative under a structure-preserving null	Established	0.338 [0.309, 0.368] vs α 0.05; matches the 3sp panel's 0.37
Filtering improves calibration under structure	Supported	0.338 to 0.057, an 83% reduction; does not reach α
FDR is controlled under mixtures	Not tested	Complete-null design only
A usable rule for choosing k exists	Not established	k remains a free, outcome-facing parameter
This transfers to empirical data	Untested here	Simulations only

8 Relationship to the empirical panel analysis

8.1 What EMMAX is actually run on

Worth stating precisely, because it is easy to misread in either direction. EMMAX here is an **ordinary per-marker test**: every marker in the panel has its own p-value, and nothing is run on eMLGs or on consensus

genotypes. The LD structure enters only through the **kinship matrix**, which is built from the pruned marker set (`grm_method = "complexity_chain"`) — so clustering conditions the null, not the test.

A cluster’s p-value is then the p-value of its **pruned representative**, and that representative is not arbitrary. `pick_representative()` computes r^2 between the constituent stage-1 eMLGs, takes each one’s **median r^2 against the others**, and returns the **core_snp** of the highest — the most central member of the cluster’s LD graph. Each stage-2 cluster has exactly one such marker (verified: $\min = \text{median} = \max = 1$ per cluster), and it is a real marker, not a synthetic genotype.

8.2 What the representative costs, measured

Two measurements on one pooled panel (V0.5_c1 nobgs env3, 294,609 SNPs in 99,787 stage-2 clusters), comparing the same EMMAX p-values at both resolutions.

Tested	tests	BH cutoff p	significant	driving QTN passing
every SNP	294,609	4.2e-04	2,467	14 of 21
LD-central representatives	99,787	1.2e-05	27	—

Fewer tests made the threshold harsher, not easier. Three times fewer tests and a BH cutoff **35 times more severe**. BH is adaptive — its cutoff is the largest p with $p \leq \alpha i/n$, so it depends on how many small p-values survive, not on n alone. Restricting to representatives removed most of the signal along with most of the tests, and the step-up never climbs.

This qualifies the rationale for filtering stated in Section 6: a lighter multiplicity correction follows only if the filter *retains* the signal. Where it does not, reducing the number of tests makes the threshold stricter and the analysis weaker on both counts.

And no driving QTN is ever its cluster’s representative — 0 of 21. That is not chance. QTN sit in unusually large clusters: the 21 QTN-bearing clusters have sizes 2, 4, 4, 5, 5, 8, 13, 21, 22, 31, 38, 55, 87, 100, 133, 141, 150, 248, 272, 311 and 354, against a genome-wide **median cluster size of 1** (mean ≈ 3). That is the expected consequence of selection extending LD around a causal site — the QTN neighbourhood *is* a large linkage block, which is what ld_w responds to — but it means the chance that any particular member is the LD-central one is about $1/\text{size}$, and for clusters of 100–350 that is negligible.

So testing representatives **never tests a QTN directly**. It relies entirely on the central marker tagging a variant that may sit at the block’s periphery, and the loss is largest exactly in the big QTN-bearing clusters where it matters most. This is the mechanism behind the cluster-level power deficits reported above, which were previously described without one.

8.3 Relationship to the empirical panel analysis

Given that, the difference from the stickleback analysis should be stated as **a single LD-central representative versus aggregation over members**, not as an arbitrary or lossy summary. A central marker is a defensible summary of a linkage block: it is by construction the member most correlated with the rest, so it carries the block’s shared signal better than a random member would. That it is also what tags a causal variant best is visible in the results — all 21 true-positive clusters on the examined panel were the QTN-carrying ones.

What remains true is narrower. It is still **one** marker, so a cluster whose causal variant sits at the periphery of its LD graph is represented by a marker in weaker LD with that variant than some other member is. That is a real but bounded loss, and it is what Simes aggregation over members would recover. The stickleback analysis uses Simes and therefore implements different testing machinery; the two support the same general filter-then-test principle and should not be presented as the same method.

9 Corrections applied after review

Defect	Resolution
“Every filtered rate sits at or below α ” was false — random reached 0.053	Restated with binomial CIs as no detectable excess over α
Mechanism example: 27 to 54 with 28 gained did not reconcile	Confirmed: one genome-wide discovery is LOST. Accounting $27 - 1 + 28 = 54$ now reported by the script
Win-rate denominator ambiguous with ties	Wins, losses and ties reported everywhere; percentages stated as non-tied
Direction reported without magnitude	Median paired change with IQR alongside every sign test; total discoveries too
80 panels pooled across heterogeneous cells and arms	Stratified by cell, arm and k ; the conclusion changed
“Licenses choosing k at all”	Withdrawn. It licenses a prespecified k only
“True discovery”	Renamed QTN-proximal throughout
“The filter must be LD-structure-based”	Weakened to the filters evaluated
“Indistinguishable at every window” conflicted with $p = 0.030$	Replaced by the stratified interaction analysis
79% rendering bug; cramped tables	Fixed; tables given explicit column widths
The cluster representative was described as an arbitrary or lossy summary	Wrong. <code>pick_representative()</code> returns the member with the highest median r^2 to the others — the LD-central marker. Restated as single-representative vs Simes aggregation, and the cost of it is now measured
“Fewer tests means a lighter multiplicity correction”	Only if the filter retains signal. Measured: 3x fewer tests gave a BH cutoff 35x MORE severe, because BH is adaptive
The <code>env_orth</code> null was dismissed as “not a true null”	REINSTATED. Its surrogates are orthogonal to the causal environment, so rejections under them are false positives. Both nulls now reported side by side
Confidence intervals were computed on 100 draws, not 1,000	A scalar named B was masked by a COLUMN named B in the input, so the grouping returned one row per input row. Intervals were ten times too wide

9.1 Open, and in progress

Re-clustering at the pilot decay settings ($n_{\text{win}} = 50$, half-decay 26 kb against the bundles’ 56 kb) is running, with $n_{\text{win}} = 10$ through the identical code path as its control. The k sweep is being extended to 40,000 and 60,000 to locate the turning point — the win rate is currently monotone up to the largest k tested, and since $k = n$ makes the two analyses identical the curve must return to ties somewhere above 20,000. Structured-null checks in the remaining cells and arms are not yet run.

10 Provenance

Generation	Maps	Half-decay	Used for
bgs5	MAP_RES 3.3e-5; 31.5% of markers share a position	55.8 kb (18 win/chr)	all cluster-level results and both manhattans
pilot_v2_v05c1	MAP_RES 3.3e-6; no shared positions	13.3 kb (95 win/chr)	marker-level power only; V0.5_c1 only
analysis_inputs	pre-bgs5	not compared	null behaviour only; V2_c1, nobgs, <code>gc_mode = none</code>

At $\rho = 0.5$ the half-decay **is** the stage-2 clustering distance, so the units themselves would differ on the newer

maps. The difference traces to decay fitting scale, not the maps: same-position marker pairs show r^2 of 0.023 against 0.022 for pairs 1 bp–10 kb apart.

Scripts: `filter_then_test_clusters.R`, `filter_then_test.R`, `filter_fdr_check.R`, `stratified_paired_table.R`, `recluster_filter_test.R`, and the three `plot_*` scripts. Parameters: α 0.05; stage-2 `min_r2_rho` 0.5, `distance_threshold` 1e5; scoring `dmax_cap` 1e5. Commits 60c8134, 88797ff, 7e595b5.

Every comparison is **paired within panel**, reported with wins, losses, ties, a median with IQR, and a sign-test p-value. Differences of means are not used.